

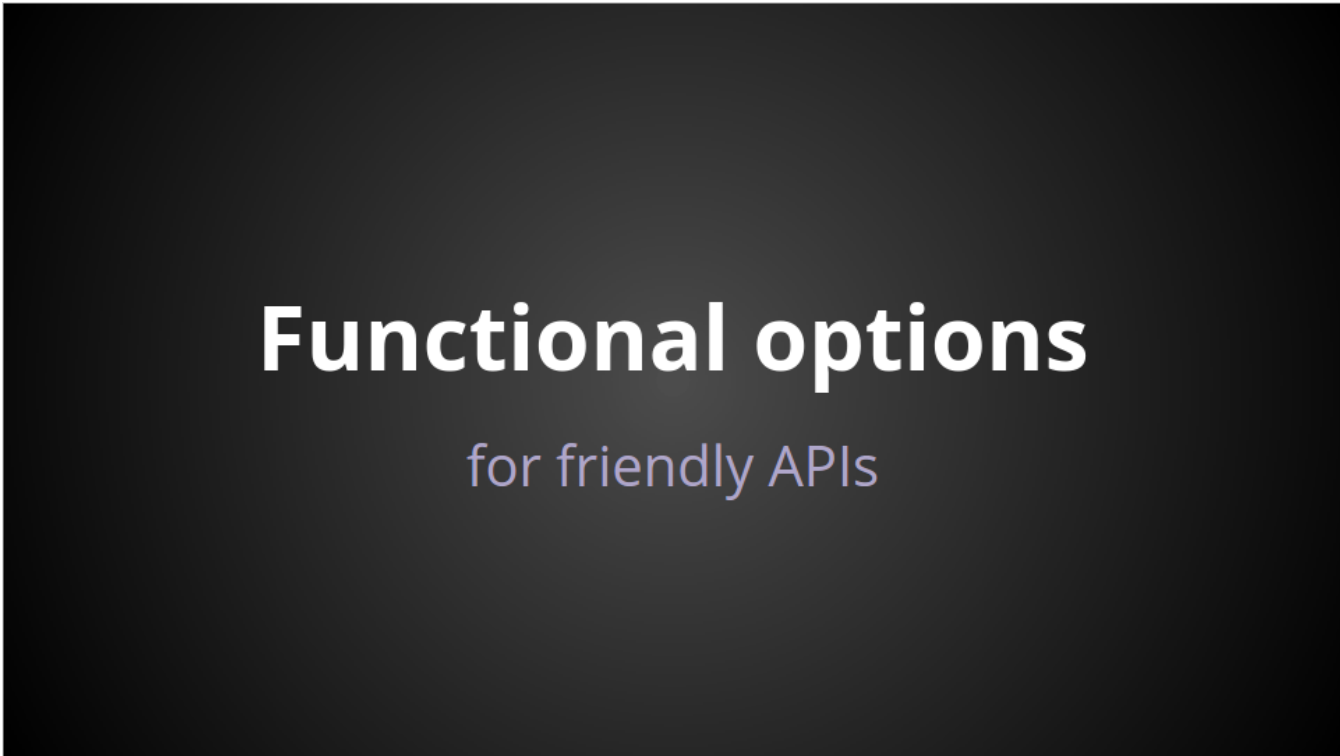
Dave Cheney

The acme of foolishness

Functional options for friendly APIs

What follows is the text of my presentation, Functional options for friendly APIs that I gave at [dotGo](#) this year. It has been edited slightly for readability.

I want to thank Kelsey Hightower, Bill Kennedy, Jeremy Saenz, and Brian Ketelsen, for their assistance in preparing this talk.



Functional options

for friendly APIs

I want to begin my talk with a story.

It is late 2014, your company is launching a revolutionary new distributed social network. Wisely, your team has chosen Go as the language for this product.

You have been tasked with writing the crucial server component. Possibly it looks a little like this.

```

package gplusplus

import "net"

type Server struct {
    listener net.Listener
}

func (s *Server) Addr() net.Addr
func (s *Server) Shutdown()

// NewServer returns a new Server listening on addr.
func NewServer(addr string) (*Server, error) {
    l, err := net.Listen("tcp", addr)
    if err != nil {
        return nil, err
    }
    srv := Server{listener: l}
    go srv.run()
    return &srv, nil
}

```

There are some unexported fields that need to be initialised, and a goroutine must be started to service incoming requests.

The package has a simple API, it is pretty easy to use.

But, there is a problem. Soon after you announce your first beta release, the feature requests start to roll in.

Features, features, features

“slow clients keep using up all my resources !!”

“does your server support TLS ?”

“can I limit the number of clients ?”

“someone is trying to DOS me, can I limit the number of connections from a single IP ?”

... and more ...

Mobile clients are often slow to respond, or stop responding altogether—you’ll need to add support for disconnecting these slow clients.

In this climate of heightened security awareness, your bug tracker starts to fill with demands to support secure connections.

Then, you get a report from a user who is running your server on a very small VPS. They need a way to limit the number of simultaneous clients.

Next is the request to rate limit concurrent connections from a group of users being targeted by a botnet.

... and on it goes.

Now, you need to change your API to incorporate all these feature requests.

An improved API, maybe

```
// NewServer returns a new Server listening on addr.  
// clientTimeout defines the maximum length of an idle  
// connection, or forever if not provided.  
// maxcons limits the number of concurrent connections.  
// maxconcurrent limits the number of concurrent  
// connections from a single IP address.  
// cert is the the TLS certificate for the connection.  
func NewServer(addr string, clientTimeout time.Duration,  
               maxcons, maxconcurrent int, cert *tls.Cert)
```

It's kind of a sign that things are not going well when the function won't easily fit on a slide.

Show of hands, who has used an API like this ?

Who has written an API like this ?

Who has had their code break while depending on an API like this ?

Obviously this solution is cumbersome and brittle. It also isn't very discoverable.

Newcomers to your package have no idea which parameters are optional, and which are mandatory.

For example, if I want to create an instance of the Server for testing, do I need to provide a real TLS certificate ? If not, what do I provide instead ?

If I don't care about `maxconns`, or `maxconcurrent` what value should I use ? Do I use zero ? Zero sounds reasonable, but depending on how the feature was implemented, that might limit you to zero total concurrent connections.

It appears to me, that writing an API like this can be easy; as long as you make it the caller's responsibility to use it correctly.

While this example could be considered an exaggeration, maliciously constructed and compounded by poor documentation, I believe that it demonstrates a real issue with ornate, brittle APIs such as this.

So now that I've defined the problem, let's look at some solutions.

Many functions make light work

```
// NewServer returns a Server listening on addr.
NewServer(addr string) (*Server, error)

// NewTLSServer returns a secure server listening on addr.
NewTLSServer(addr string, cert *tls.Cert) (*Server, error)

// NewServerWithTimeout returns a Server listening on addr that
// disconnects idle clients.
NewServerWithTimeout(addr string, timeout time.Duration) (*Server, error)

// NewTLSServerWithTimeout returns a secure Server listening
// on addr that disconnects idle clients.
NewTLSServerWithTimeout(addr string, cert *tls.Cert,
                        timeout time.Duration) (*Server, error)
```

Rather than trying to provide one single function which must cater for every permutation, a solution might be to create a set of functions.

With this approach, when callers need a secure server they can call the TLS variant.

When they need to establish a maximum duration for idle connections, they can use the variant that takes a timeout.

Unfortunately, as you can see, providing every possible permutation can quickly become overwhelming.

Let's move on to others way of making your API configurable.

Passing a configuration struct

```
// A Config structure is used to configure the Server.
type Config struct {
    // Timeout sets the amount of time before closing
    // idle connections, or forever if not provided.
    Timeout time.Duration

    // The server will accept TLS connections is the
    // certificate provided.
    Cert *tls.Cert
}

func NewServer(addr string, config Config) (*Server, error)
```

A very common solution is to use a configuration struct.

This has some advantages.

Using this approach, the configuration struct can grow over time as new options are added, while the public API for creating a server itself remains unchanged.

This method can lead to better documentation.

What was once a massive comment block on the `NewServer` function, becomes a nicely documented struct.

Potentially it also enables the callers to use the zero value to signify they they want the default behaviour for a particular configuration option.

Limited options

```
type Config struct {  
    // The port to listen on, if unset defaults to 8080.  
    Port int  
}  
  
func main() {  
    srv, _ := NewServer("localhost", Config{  
        Port: 0, // opps, can't do this  
    },  
    )  
}
```

However, this pattern is not perfect.

It has trouble with defaults, especially if the zero value has a well understood meaning.

For example, in the config structure shown here, when Port is not provided, NewServer will return a *Server for listening on port 8080.

But this has the downside that you can no longer explicitly set Port to 0 and have the operating system automatically choose a free port, because that explicit 0 is indistinguishable from the fields' zero value.

"I just want a server, I don't want to have to think about it"

```
func NewServer(addr string, config Config) (*Server, error)  
  
func main() {  
    srv, _ := NewServer(addr, Config{}) // because reasons  
    ...  
}
```

Most of the time, users of your API will be expecting to use its default behaviour.

Even though they do not intend to change any of the configuration parameters, those callers are still required to pass *something* for that second argument.

So, when people read your tests or your example code, trying to figure out how to use your package, they'll see this magic empty value, and it'll become enshrined in the collective unconsciousness.

[and] to me, this just feel wrong.

Why should users of your API be required to construct an empty value, simply to satisfy the signature of the function ?

Maybe adding indirection will help

```
func NewServer(addr string, config *Config) (*Server, error)

func main() {
    srv, _ := NewServer("localhost", nil) // accept the defaults

    config := Config{Port: 9000}
    srv2, _ := NewServer("localhost", &config)

    config.Port = 9001 // what happens now ?
    ...
}
```

A common solution to this empty value problem is to pass a pointer to the value instead, thereby enabling callers to use `nil` rather than constructing an empty value.

In my opinion this pattern has all the problems of the previous example, and it adds a few more.

We still have to pass *something* for this function's second argument, but now this value could be `nil`, and most of the time *will* be `nil` for those wanting the default behaviour.

It raises the question, is there a difference between passing `nil`, and passing a pointer to an empty value ?

More concerning to both the package's author, and its callers, is the `Server` and the caller can now share a reference to the same configuration value. Which gives rise to questions of what happens if this value is mutated after being passed to the `NewServer` function ?

I believe that well written APIs should not require callers to create dummy values to satisfy those rarer use cases.

I believe that we, as Go programmers, should work hard to ensure that nil is never a parameter that needs to be passed to any public function.

And when we do want to pass configuration information, it should be as self explanatory and as expressive as possible.

So now with these points in mind, I want to talk about what I believe are some better solutions.

Variadic configuration

```
func NewServer(addr string, config ...Config) (*Server, error)

func main() {
    srv, _ := NewServer("localhost") // defaults

    // timeout after 5 minutes, 10 clients max
    srv2, _ := NewServer("localhost", Config{
        Timeout: 300 * time.Second,
        MaxConns: 10,
    })
    ...
}
```

To remove the problem of that mandatory, yet frequently unused, configuration value, we can change the NewServer function to accept a variable number of arguments.

Instead of passing nil, or some zero value, as a signal that you want the defaults, the variadic nature of the function means you don't need to pass *anything at all*.

And in my book this solves two big problems.

First, the invocation for the default behaviour becomes as concise as possible.

Secondly, NewServer now only accepts Config values, not pointers to config values, removing nil as a possible argument, and ensuring that the caller cannot retain a reference to the server's internal configuration.

I think this is a big improvement.

But if we're being pedantic, it still has a few problems.

Obviously the expectation is for you to provide at most one Config value. But as the function signature is variadic, the implementation has to be written to cope with a caller passing multiple, possibly contradictory, configuration structs.

Is there a way to use a variadic function signature *and* improve the expressiveness of configuration parameters when needed ?

I think that there is.

Functional options

```
func NewServer(addr string, options ...func(*Server)) (*Server, error)

func main() {
    srv, _ := NewServer("localhost") // defaults

    timeout := func(srv *Server) {
        srv.timeout = 60 * time.Second
    }

    tls := func(srv *Server) {
        config := loadTLSConfig()
        srv.listener = tls.NewListener(srv.listener, &config)
    }

    // listen securely with a 60 second timeout
    srv2, _ := NewServer("localhost", timeout, tls)
}
```

At this point I want to make it clear that the idea of functional options comes from a blog post titled. [Self referential functions and design](#) by Rob Pike, published in January this year. I encourage everyone here to read it.

The key difference from the previous example, and in fact all the examples so far, is customisation of the Server is performed not with configuration parameters stored in a structure, but with *functions* which operate on the Server value itself.

As before, the variadic nature of the function's signature gives us the compact behaviour for the default case.

When configuration is required, I pass to NewServer functions which operate on the Server value as an argument.

The timeout function simply changes the timeout field of any *Server value passed to it.

The tls function is a little more complicated. It takes a *Server value and wraps the original listener value inside a tls.Listener, thereby transforming it into a secure listener.

Functional options, cont

```
// NewServer returns a Server listening on addr.
func NewServer(addr string, options ...func(*Server)) (*Server, error) {
    l, err := net.Listen("tcp", addr)
    if err != nil {
        return nil, err
    }

    srv := Server{listener: l}

    for _, option := range options {
        option(&srv)
    }

    return &srv, nil
}
```

Inside NewServer, applying these options is straightforward.

After opening a net.Listener, we declare a Server instance using that listener.

Then, for each option function provided to NewServer, we call that function, passing in a pointer to the Server value that was just declared.

Obviously, if no option functions were provided, there is no work to do in this loop and so srv is unchanged.

And that's all there is too it.

Using this pattern we can make an API that has

- sensible defaults
- is highly configurable
- can grow over time
- self documenting
- safe for newcomers
- and never requires nil or an empty value to keep the compiler happy

In the few minutes I have remaining I'd like to show you how I improved one of my own packages by converting it to use functional options.

Term (original)

```
package main

import "github.com/pkg/term"

func main() {
    t, err := term.Open("/dev/ttyUSB0")
    // handle error
    err = t.SetSpeed(115200)
    // handle error
    err = t.SetRawMode()
    // handle error
    ...
}
```

I'm an amateur hardware hacker, and many of the devices I work with use a USB serial interface. A so a few months ago I wrote a [terminal handling package](#).

In the prior version of this package, to open a serial device, change the speed and set the terminal to raw mode, you'd have to do each of these steps individually, checking the error at every stage.

Even though this package is trying to provide a friendlier interface on an even lower level interface, it still left too many procedural warts for the user.

Let's take a look at the package after applying the functional options pattern.

Term (improved)

```
package main

import "github.com/pkg/term"

func main() {
    // just open the terminal
    t, _ := term.Open("/dev/ttyUSB0")

    // open at 115200 baud in raw mode
    t2, _ := term.Open("/dev/ttyUSB0", term.Speed(115200),
term.RawMode)
    ...
}
```

By converting the `Open` function to use a variadic parameter of function values, we get a much cleaner API.

In fact, it's not just the `Open` API that improves, the grind of setting an option, checking an error, setting the next option, checking the error, that is gone as well.

The default case, still just takes one argument, the name of the device.

For more complicated use cases, configuration functions, defined in the `term` package, are passed to the `Open` function and are applied in order before returning.

This is the same pattern we saw in the previous example, the only thing that is different is rather than being anonymous, these are public functions. In all other respects their operation is identical.

We'll take a look at how `Speed`, `RawMode`, and `Open`, are implemented on the next slide.

term.Open

```
package term

// RawMode places the terminal into raw mode.
func RawMode(t *Term) error { return t.setRawMode() }

// Speed sets the baud rate option for the terminal.
func Speed(baud int) func(*Term) error {
    return func(t *Term) error {
        return t.setSpeed(baud)
    }
}

// Open opens an asynchronous communications port.
func Open(dev string, options ...func(*Term) error) (*Term, error) {
    t, err := openTerm(dev)
    // handle err
    for _, option := range options {
        err := option(t)
        // handle err, cleanup
    }
    return t, nil
}
```

`RawMode` is the easiest to explain. It just a function whose signature is compatible with `Open`.

Because `RawMode` is declared in the same package as `Term`, it can access the private fields and call private methods declared on the `Term` type, in this case calling the private `setRawMode` helper.

`Speed` is also just a regular function, however it does not match the signature `Open` requires. This is because `Speed` itself requires an argument; the baud rate.

`Speed` returns an anonymous function which is compatible with the `Open` function's signature, which closes over the baud rate parameter, capturing it for later when the function is applied.

Inside the call to `Open`, we first open the terminal device with the `openTerm` helper.

Next, just as before, we range over the slice of options functions, calling each one in turn passing in `t`, the pointer to our `term`. `Term` value.

If there is an error applying any function then we stop at that point, clean up and return the error to the caller.

Otherwise, returning from the function, we've now created and configured a `Term` value to the caller's specifications.

Functional options

- ✓ let you write beautiful APIs that can grow over time
- ✓ enable the default use case to be its simplest
- ✓ give more readable, meaningful parameters
- ✓ provide direct control over the initialisation of complex values

In summary

- Functional options let you write APIs that can grow over time.
- They enable the default use case to be the simplest.
- They provide meaningful configuration parameters.
- Finally they give you access to the entire power of the language to initialize complex values.

In this talk, I have presented many of the existing configuration patterns, those considered idiomatic and commonly in use today, and at every stage asked questions like:

- Can this be made simpler ?
- Is that parameter necessary ?
- Does the signature of this function make it easy for it to be used safely ?
- Does the API contain traps or confusing misdirection that will frustrate ?

I hope I have inspired you to do the same. To revisit code that you have written in the past and pose yourself these same questions and thereby improve it.

Thank You!

@davecheney

Thank you.

Related Posts:

1. [How to include C code in your Go package](#)
2. [Do not fear first class functions](#)
3. [Using go test, build and install](#)
4. [Simple profiling package moved, updated](#)

This entry was posted in Go, Programming and tagged dotgo, functional options on October 17, 2014
[<https://dave.cheney.net/2014/10/17/functional-options-for-friendly-apis>] .
