

[Ewan Higgs](#) [home](#)

I consider myself mainly a Python developer these days, but coming from a history of C++ development, I've been really interested in seeing how Rust pans out. Could it be an ideal language to use when I need the performance that Python just can't provide? Can I happily replace C or C++ (using boost) with Rust?

I set out to see the current state of binding Python to Rust.

Python

Let's begin with a fibonacci sequence calculator. It's pretty basic in Python:

```
1  def fib(n):
2      if n < 2:
3          return 1
4      prev1 = 1
5      prev2 = 1
6      for i in range(1, n):
7          next = prev1 + prev2
8          prev2 = prev1
9          prev1 = next
10
11     return prev1
```

python_example.py hosted with ❤ by [GitHub](#)

[view raw](#)

C

C has a lot of boiler plate where I could just use the `ffi` module. But using the actual API is worth it in case you want to do any sort of error handling in the C side. For example, if you pass a wrong type through `ffi`, you basically just segfault. You can wrap your C code with error handling in Python. But if the intent is to make a function you call in a loop, you want the error checking to also be fast.

If I'm honest, I always feel like I'm putting in a lot of bugs when I use the C API. Especially when I compare it with the Cython output.

```
1  #include <Python.h>
2  #include <stdint.h>
3
4  static uint64_t _fib(uint64_t n) {
5      if (n == 1 || n == 2) {
6          return 1;
7      }
8      uint64_t prev1 = 1;
9      uint64_t prev2 = 1;
10     for (uint64_t i = 1; i < n; ++i) {
11         uint64_t next = prev1 + prev2;
```

```

12     prev2 = prev1;
13     prev1 = next;
14 }
15 return prev1;
16 }
17
18 static PyObject* fib(PyObject* self, PyObject* args) {
19     int64_t n;
20
21     if (!PyArg_ParseTuple(args, "L", &n)) {
22         return NULL;
23     }
24     if (n < 0) {
25         return NULL;
26     }
27     uint64_t f = _fib(n);
28     return Py_BuildValue("L", f);
29 }
30
31 static PyMethodDef c_python_methods[] =
32 {
33     {"fib", fib, METH_VARARGS, "Compute a fibonacci sequence value."},
34     {NULL, NULL, 0, NULL}
35 };
36
37 static struct PyModuleDef c_python_module = {
38     PyModuleDef_HEAD_INIT,
39     "c_python_exaplme",
40     "Example module",
41     -1,
42     c_python_methods,
43     NULL,
44     NULL,
45     NULL,
46     NULL
47 };
48
49 PyMODINIT_FUNC
50 PyInit_c_python_example(void)
51 {
52     return PyModule_Create(&c_python_module);
53 }

```

c_python_example.c hosted with ❤ by [GitHub](#)

[view raw](#)

```

1 .PHONY: all
2
3 CFLAGS=$(shell pkg-config --libs --cflags python3) -O3 -Wall -Werror -pedantic -march=native
4 all: c_python_example.so

```

```

5
6 c_python_example.so: c_python_example.c
7     $(CC) $(CFLAGS) $< -o $@
8
9 clean:
10     rm c_python_example.so

```

Makefile hosted with ♥ by GitHub

[view raw](#)

C++ (Boost.Python)

Here's what the equivalent looks like in [Boost.Python](#). Boost.Python is a really nice library for wrapping C++ to call from Python. Except if you manage to make a mistake you get all the associated C++ compilation errors some of us love to hate. It's allegedly gotten better, but in making this example, I still found some error explosions which take some searching around on the internet to figure out.

Sadly Boost.Python only seems to support Python 2. I know a lot of people using Python 2 who have no interest in migrating to Python 3, but if you end up wrapping a lot of your code in Boost.Python, you're really sealing your fate. I don't recommend it.

```

1 #include <boost/python.hpp>
2
3 static uint64_t fib(uint64_t n) {
4     if (n == 1 || n == 2) {
5         return 1;
6     }
7
8     uint64_t prev1 = 1;
9     uint64_t prev2 = 1;
10    for (uint64_t i = 1; i < n; ++i) {
11        uint64_t next = prev1 + prev2;
12        prev2 = prev1;
13        prev1 = next;
14    }
15    return prev1;
16 }
17
18 BOOST_PYTHON_MODULE(cc_python_example) {
19     using namespace boost::python;
20     def("fib", fib);
21 }

```

cc_python_example.cc hosted with ♥ by GitHub

[view raw](#)

```

1 .PHONY: all
2
3 CFLAGS=$(shell pkg-config --libs --cflags python) -lboost_python -O3 -Wall -Werror -pedantic

```

```

4  all: cc_python_example.so
5
6  cc_python_example.so: cc_python_example.cc
7      $(CXX) $(CFLAGS) $< -o $@
8
9  clean:
10      rm cc_python_example.so

```

Makefile hosted with ❤ by GitHub

[view raw](#)

Cython

Cython is a fantastic tool for writing Python bindings. It's a dialect of Python that looks so much like Python that you can almost copy paste your code into a `.pyx` module, configure `setuptools` to find it, and you're running at a much faster speed. If you take a look at the Cython code, set the types on the variables, turn off array bounds checking, and other tweaks, it competes with hand written C. It's really awesome.

```

1  def fib(int n):
2      if n < 2:
3          return 1
4      cdef unsigned long long prev1 = 1
5      cdef unsigned long long prev2 = 1
6      cdef int i = 1
7      while i < n:
8          next = prev1 + prev2
9          prev2 = prev1
10         prev1 = next
11         i += 1
12     return prev1

```

cython_python_example.pyx hosted with ❤ by GitHub

[view raw](#)

Rust (rust-cpython)

The best library I've found for writing modules in Rust is [rust-cpython](#). Using it took some trial and error since it's so new and I wouldn't call myself a Rustacean just yet. But it was actually a very pleasant experience. With a little manipulation of the `Cargo.toml` file (used to define a reproducible build in Rust), you can build against Python 2 or 3. None of the other wrappers, aside from Cython can do this.

Finally, here's how the Fibonacci number generator looks in Rust:

```

1  [package]
2  name = "rust_python_example"
3  version = "0.1.0"
4  authors = []

```

```

5
6 [lib]
7 name = "rust_python_example"
8 crate-type = ["dylib"]
9
10 [dependencies]
11 interpolate_idents = "*"
12
13 [dependencies.cpython]
14 version = "*"
15 default-features = false
16 features = ["python3-sys"]

```

Cargo.toml hosted with ♥ by GitHub

[view raw](#)

```

1  #![feature(plugin)]
2  #![plugin(interpolate_idents)]
3
4  #[macro_use] extern crate cpython;
5
6  use cpython::{PyResult, Python, PyTuple, PyErr, exc, ToPyObject, PyObject};
7
8  mod fib {
9
10     pub fn fib(n : u64) -> u64 {
11         if n < 2 {
12             return 1
13         }
14         let mut prev1 = 1;
15         let mut prev2 = 1;
16         for _ in 1..n {
17             let new = prev1 + prev2;
18             prev2 = prev1;
19             prev1 = new;
20         }
21         prev1
22     }
23 }
24
25 py_module_initializer!(librust_python_example, |_py, m| {
26     try!(m.add("__doc__", "Module documentation string"));
27     try!(m.add("fib", py_fn!(fib)));
28     Ok(())
29 });
30
31 fn fib<'p>(py: Python<'p>, args: &PyTuple<'p>) -> PyResult<'p, u64> {
32     let arg0 = match args.get_item(0).extract::<u64>() {
33         Ok(x) => x,
34         Err(_) => {

```

```
35         let msg = "Fib takes a number greater than 0";
36         let pyerr = PyErr::new_lazy_init(py.get_type::<exc::ValueError>(), Some(msg.to_py));
37         return Err(pyerr);
38     }
39 };
40 Ok(fib::fib(arg0))
41 }
```

lib.rs hosted with ❤ by GitHub [view raw](#)

I ran some basic tests so see how the performance fares, and it turns out that it's pretty competitive. I wouldn't put too much stock into the actual numbers aside from the fact that they're within an order of magnitude of each other. I used GCC for C and Rust is built on LLVM.

The test is a really noddy one. It's just using 'timeit' on the 91st fibonacci number from the IPython shell:

```
timeit fib(90)
```

Python Version	Implementing language	timeit result
Python 3.4.2	Python	100000 loops, best of 3: 7.05 µs per loop
Python 3.4.2	C	1000000 loops, best of 3: 288 ns per loop
Python 3.4.2	Rust	1000000 loops, best of 3: 229 ns per loop
Python 3.4.2	Cython	10000000 loops, best of 3: 192 ns per loop
Python 2.7.10	Python	100000 loops, best of 3: 6.11 µs per loop
Python 2.7.10	C++	1000000 loops, best of 3: 219 ns per loop
Python 2.7.10	Rust	1000000 loops, best of 3: 177 ns per loop
Python 2.7.10	Cython	10000000 loops, best of 3: 171 ns per loop

As it's a microbenchmark, the only thing I would say with any amount of confidence is that the Rust API wrapper probably isn't getting in the way too much.

Current work

Some Rustaceans/Pythonistas are working on this infrastructure and it's pretty exciting. Other pieces of work being put into place are integration with [Python's setuptools](#) so you can run `python setup.py install` and it will be able to build the Python modules just like you would expect.

Current issues

1. I haven't found anyone who is looking to support Numpy yet, so anyone who wants to do their numeric/scientific development in Python/Rust should hold off at the moment.
2. I don't see how to specify the library name in Cargo. This means all the shared objects are always prefixed with `lib`. This means you need to import the module as `import libXYZ` which isn't really what you want.

Further work

I haven't tested managing objects on the Rust side, but this is apparently a place where Python can potentially get some performance improvements. For example, [collections.OrderedDict was implemented in raw Python until Python 3.5](#). Of course, the ability to use [RAII](#) on the Rust side should make writing containers in Rust a lot easier than in C.

Conclusion

Rust looks like a really great language to implement Python modules. It's not there yet if you want to integrate with the numeric/stats/machine learning libraries that Python offers. But with some love and more community effort, I could see a community building in this direction. I think it's a very promising start and I hope it continues.

Related Posts

Ewan Higgs

github.com/ehiggs

bitbucket.org/ewanhiggs

be.linkedin.com/in/ewanhiggs